

KISHORE A

Bengaluru, Karnataka

Phone: +91 6382707597 | Mail: a.kishore.dev@gmail.com | LinkedIn: [linkedin.com/in/kishore-ai](https://www.linkedin.com/in/kishore-ai)

GitHub: github.com/a-kishore-dev | Portfolio: kishoreai.me

Professional Summary

AI/ML Engineer with hands-on experience building production-grade RAG systems, fine-tuning large language models with QLoRA, and shipping end-to-end machine learning pipelines from feature engineering to deployment. Implemented transformer and GPT architectures from scratch in PyTorch to ground applied work in first-principles understanding of LLM internals. Seeking an AI Engineer, ML Engineer, or GenAI role to design and ship production AI systems.

Education

Panimalar Engineering College Chennai City Campus

Nov 2022 – May 2026

Bachelor of Technology in Artificial Intelligence and Data Science

Chennai, Tamil Nadu

CGPA: 8.67

Technical Skills

Languages: Python, SQL

LLM & Generative AI: LangChain, LlamaIndex, RAGAS, HyDE, FAISS, ChromaDB, Qdrant, QLoRA/LoRA, Prompt Engineering

Deep Learning & ML: PyTorch, TensorFlow, Scikit-learn, XGBoost, LightGBM, CatBoost, SHAP

MLOps & Tools: MLflow, LangSmith, Weights & Biases, Git, Jupyter

Cloud & Deployment: AWS (EC2, S3, Lambda, SageMaker), Flask, Streamlit, Vercel, Render

Projects

Production RAG System | *Python, LlamaIndex, Qdrant, BM25, RAGAS, Gemini*

Live | GitHub

- Built a production-grade RAG pipeline combining BM25 keyword search with dense vector retrieval, HyDE query enhancement, and CrossEncoder semantic reranking, deployed on Streamlit Cloud for public access.
- Evaluated retrieval and generation quality against a self-authored ground-truth QA set using the RAGAS framework, achieving **1.0 Context Recall**, **0.9722 Context Precision**, and **0.9455 Faithfulness**.

Llama-3.2-3B Text-to-SQL Fine-Tuning | *Python, PyTorch, Unsloth, PEFT, QLoRA*

HuggingFace | GitHub

- Fine-tuned Llama 3.2 3B for text-to-SQL generation via QLoRA, training only **0.75% of 3.21B parameters** using 4-bit quantization and LoRA adapters on a free-tier Colab T4.
- Improved **ROUGE-1 score from 0.29 to 0.85 (+192%)** over the base model on 200 held-out test examples; published the merged model on HuggingFace Hub.

Loan Default Risk Predictor | *Python, LightGBM, XGBoost, SHAP, MLflow*

GitHub

- Engineered **40+ features** from **307,511 loan records** across 5 relational tables, aggregating credit bureau history, payment status, and prior application data.
- Benchmarked 6 models on an 11.4% class-imbalanced dataset using scale_pos_weight tuning; LightGBM achieved the best **validation ROC-AUC of 0.7662**, tracked via MLflow experiment logging.

Experience

Infosys Springboard

Oct 2024 – Nov 2024

AI Intern

Remote

- Trained a Random Forest classifier on a food allergen dataset, achieving a **99% F1-score** across nut, gluten, and dairy categories.
- Deployed a Flask REST API on Render with **sub-second inference latency**, paired with a Streamlit interface for real-time allergen prediction.

Certifications

Deep Learning Specialization – DeepLearning.AI (2024)

Machine Learning Specialization – DeepLearning.AI (2024)

Natural Language Processing Specialization – DeepLearning.AI (2025)

AWS Academy Cloud Foundations – AWS Academy Graduate

Achievements

- MozoHack 6.0 Hackathon – Top 10 Finalist** (out of 250+ teams), SRMKZILLA, SRM University (2025) – Developed the blink-detection pipeline and Gemini LLM chat integration for NeuroBlink, an assistive eye-blink-to-Morse-code communication system.